www.locuz.com

Intel oneAPI Series - Deep Dive (Live Demo and Hands-on): oneAPI and DPC++ on Intel DevCloud

Presented By: Mandeep Kumar

Converge to the Cloud

Agenda

- Overview of Intel oneAPI and Data Parallel C++
- Introduction to the Intel DevCloud
- Setup of Intel DevCloud and JupyterLab Environment
- DPC++ Program Structure
- Demonstration of Intel VTune Profiler on Intel DevCloud
- Demonstration of Intel DPC++ Compatibility Tool on Intel DevCloud

Learning Objectives

- Articulate how oneAPI can help to solve the challenges of programming in a heterogeneous world
- Understand the DPC++ language and programming model
- Profile a DPC++ application using Intel VTune Profiler on Intel DevCloud
- Learn how to migrate CUDA code to Data Parallel C++ using the Intel DPC++ Compatibility tool

Programming Challenges for Multiple Architectures

- Growth in specialized workloads
- Variety of data-centric hardware required
- Separate programming models and toolchains for each architecture are required today
- Software development complexity limits freedom of architectural choice

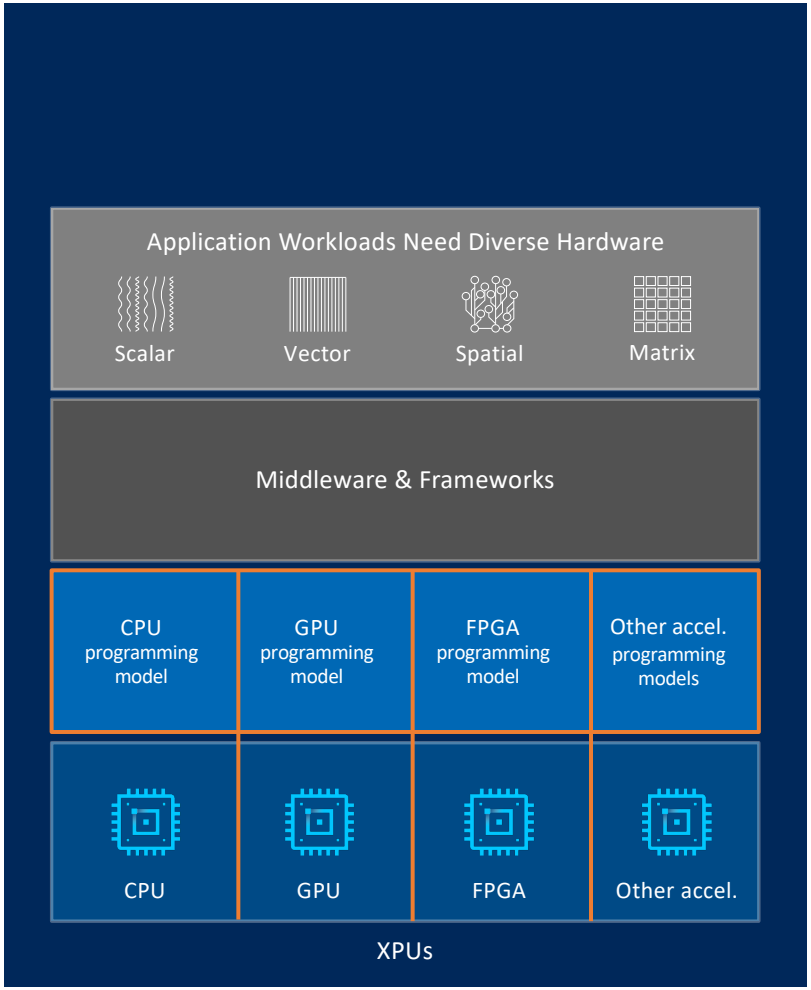

Introducing oneAPI

- Cross-architecture programming that delivers freedom to choose the best hardware
- Based on industry standards and open specifications
- Exposes cutting-edge performance features of latest hardware
- Compatible with existing high-performance languages and programming models including C++, OpenMP, Fortran, and MPI

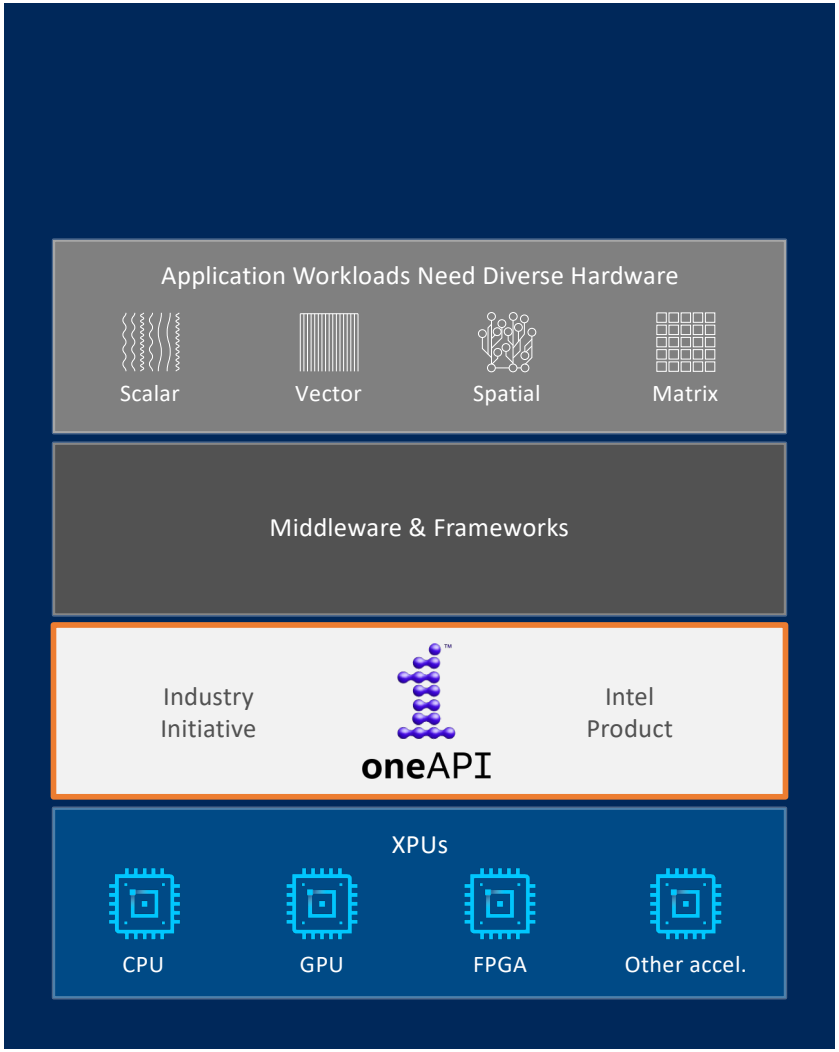

oneAPI Industry Initiative: Break the Chains of Proprietary Lock-in

- A cross-architecture language based on C++ and SYCL standards
- Powerful libraries designed for acceleration of domain-specific functions
- Low-level hardware abstraction layer
- **Open to promote community and industry collaboration**
- **Enables code reuse across architectures and vendors**

oneAPI

The productive, smart path to freedom for accelerated computing from the economic and technical burdens of proprietary programming models

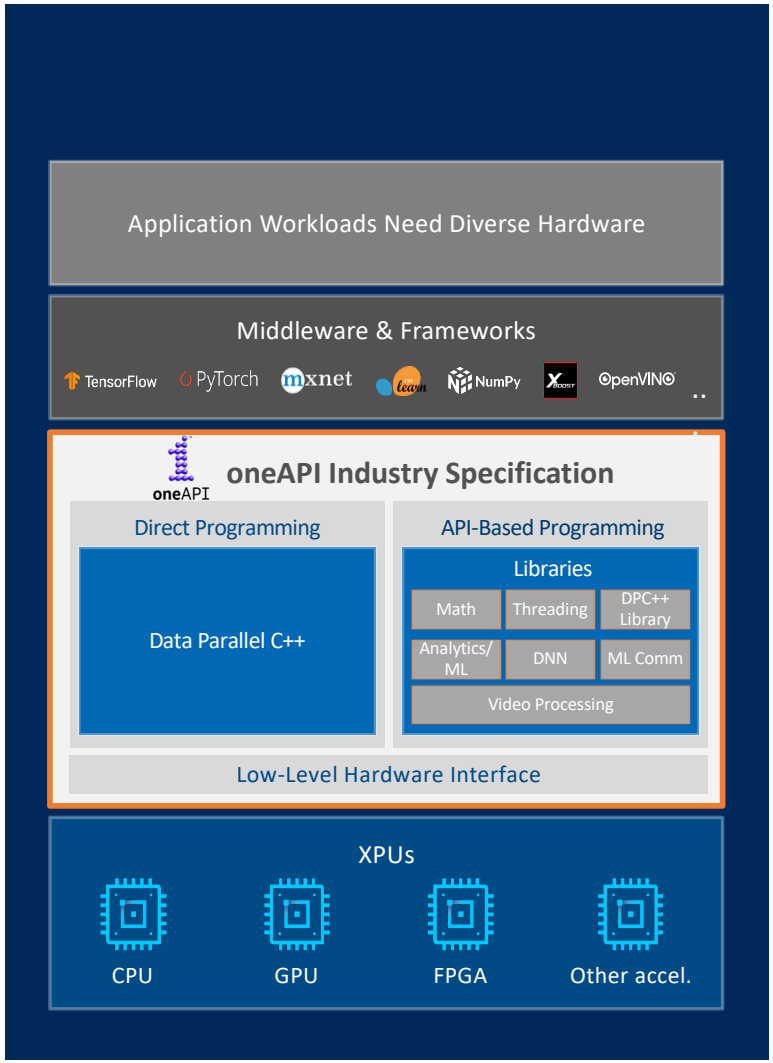

visit oneapi.com for more details

Data Parallel C++: Standards-based, Cross-architecture Language

Parallelism, productivity and performance for CPUs and Accelerators

- Delivers accelerated computing by exposing hardware features
- Allows code reuse across hardware targets, while permitting custom tuning for specific accelerators
- Provides an open, cross-industry solution to single architecture proprietary lock-in

Based on C++ and SYCL

- Delivers C++ productivity benefits, using common, familiar C and C++ constructs
- Incorporates SYCL from the Khronos Group to support data parallelism and heterogeneous programming

Community Project to drive language enhancements

- Provides extensions to simplify data parallel programming
- Continues evolution through open and cooperative development

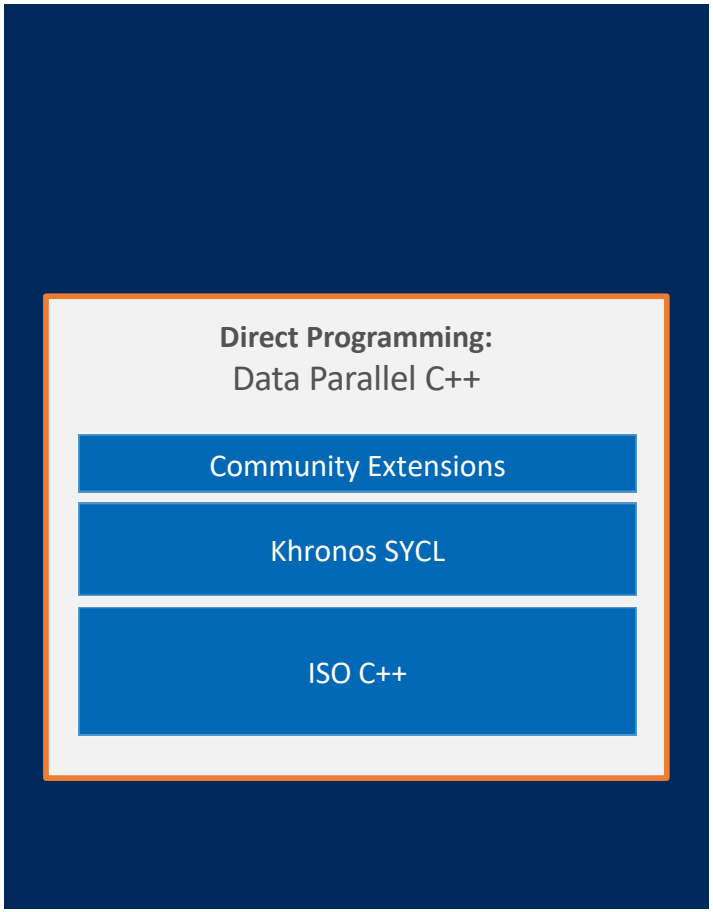

Apply your skills to the next innovation, not rewriting software for the next hardware platform

The open source and Intel DPC++/C++ compiler supports Intel CPUs, GPUs, and FPGAs. Codeplay announced a DPC++ complier that targets Nvidia GPUs.

Intel oneAPI Toolkits

A complete set of proven developer tools expanded from CPU to XPU

Intel® oneAPI Base Toolkit

Native Code Developers

A core set of high-performance tools for building C++, Data Parallel C++ applications & oneAPI library-based applications

Add-on Domain-specific Toolkits

Specialized Workloads

Intel® oneAPI Tools for HPC

Deliver fast Fortran, OpenMP & MPI applications that scale

Intel® oneAPI Tools for IoT

Build efficient, reliable solutions that run at network's edge

Intel® oneAPI Rendering Toolkit

Create performant, high-fidelity visualization applications

Toolkits powered by oneAPI

Data Scientists & AI Developers

Intel® AI Analytics Toolkit

Accelerate machine learning & data science pipelines with optimized DL frameworks & high-performing Python libraries

Intel® Distribution of OpenVINO™ Toolkit

Deploy high performance inference & applications from edge to cloud

Transition from Parallel Studio XE with no Disruption

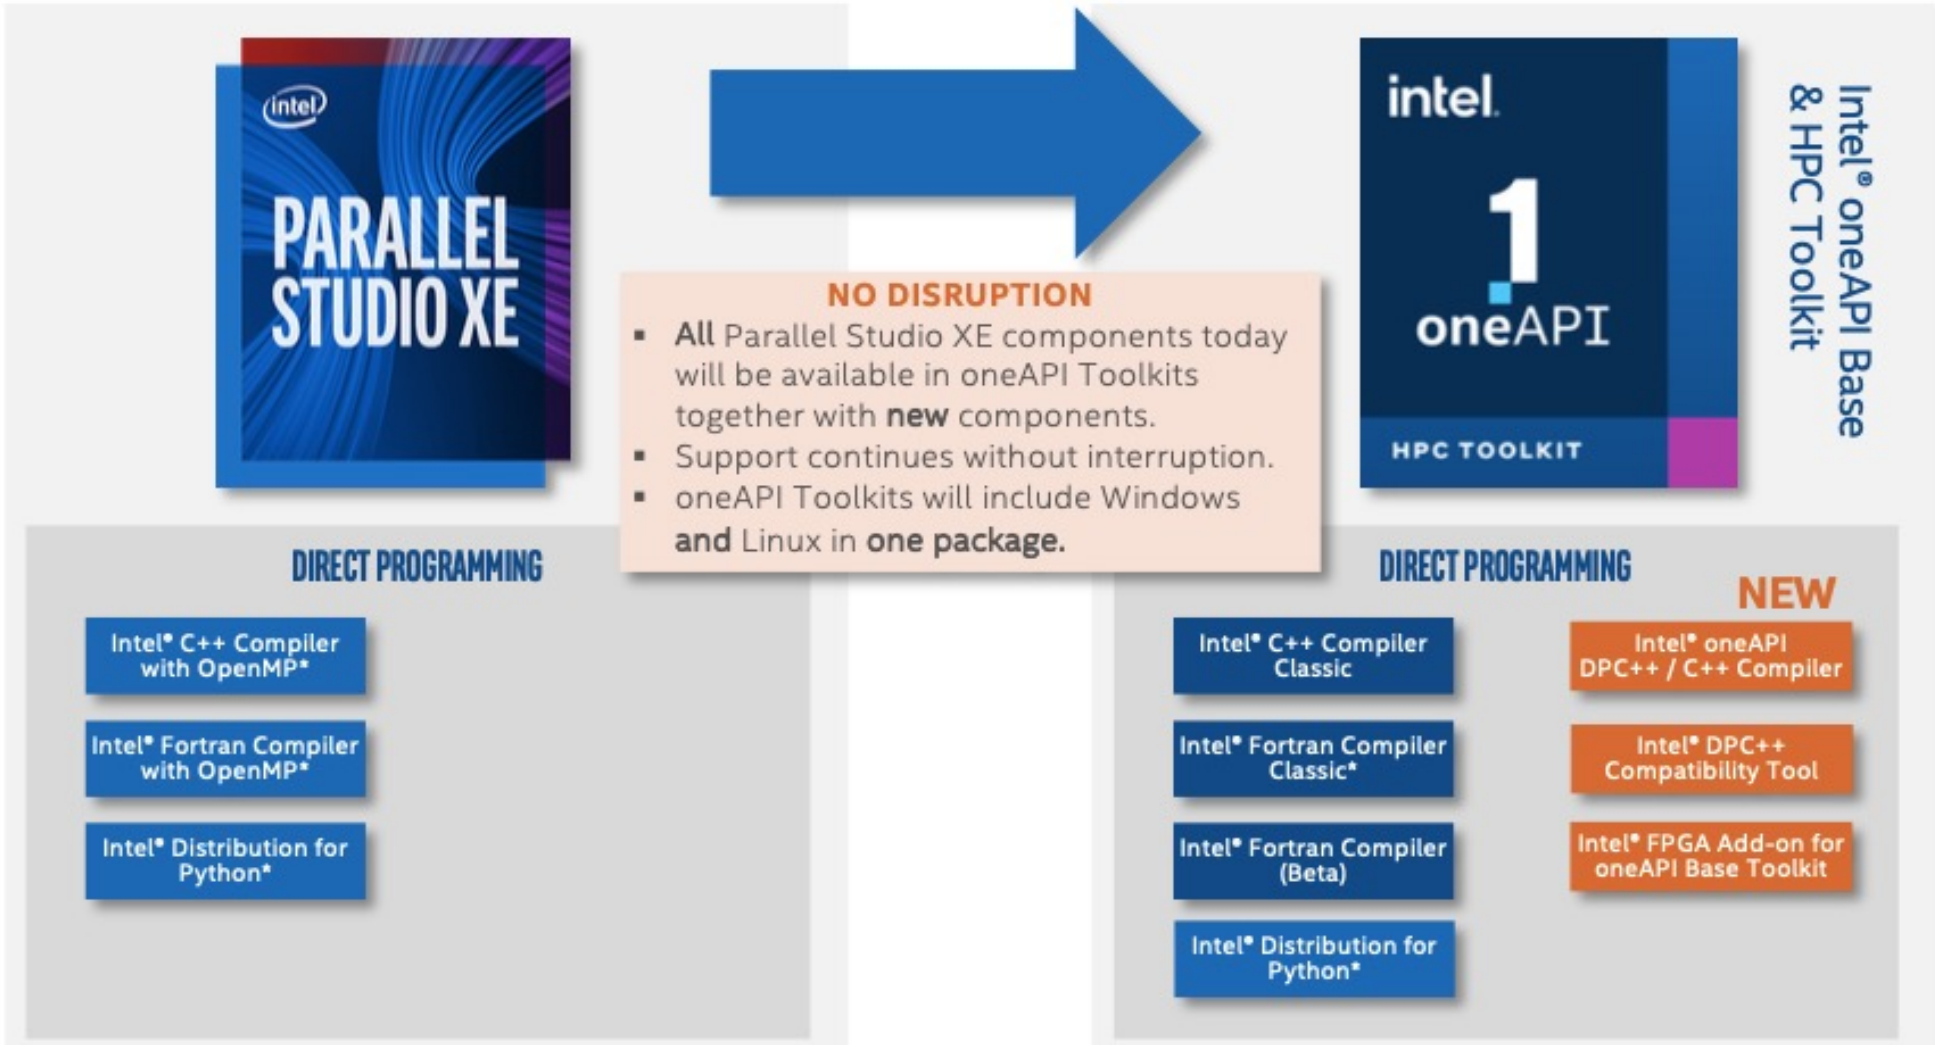

Transition from Parallel Studio XE with no Disruption

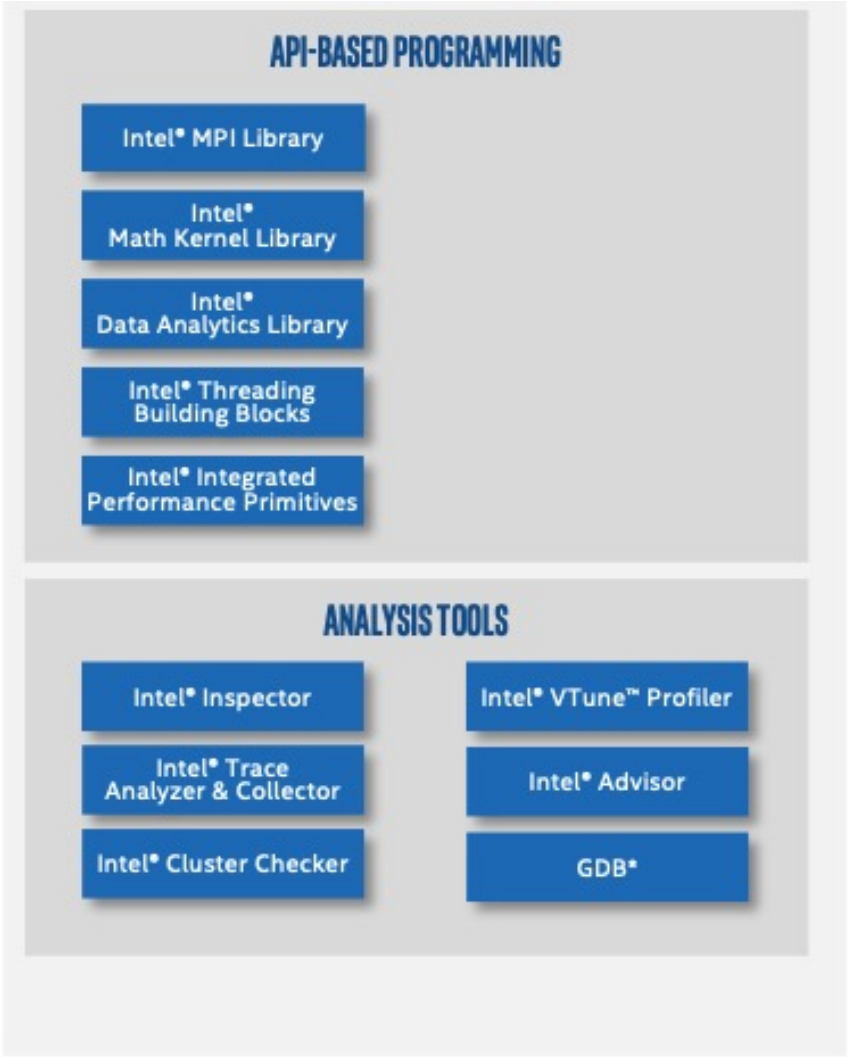

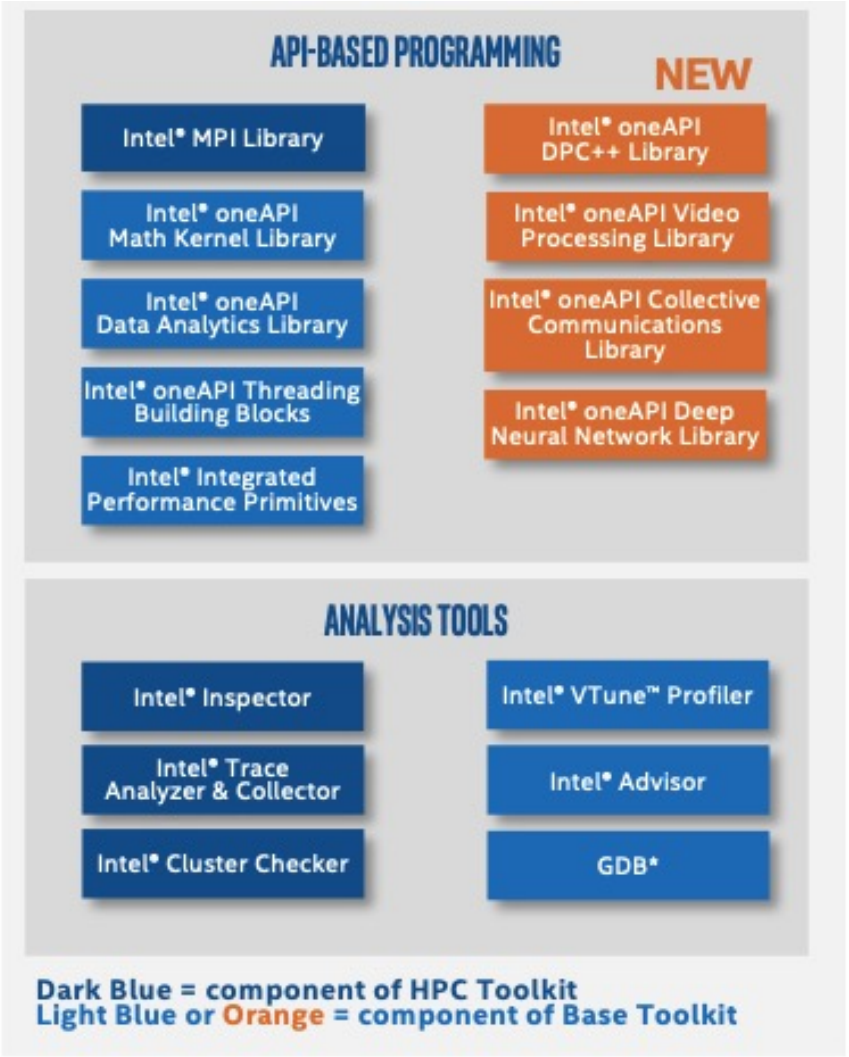

Transition from Parallel Studio XE with no Disruption

Transitioning to new Editions

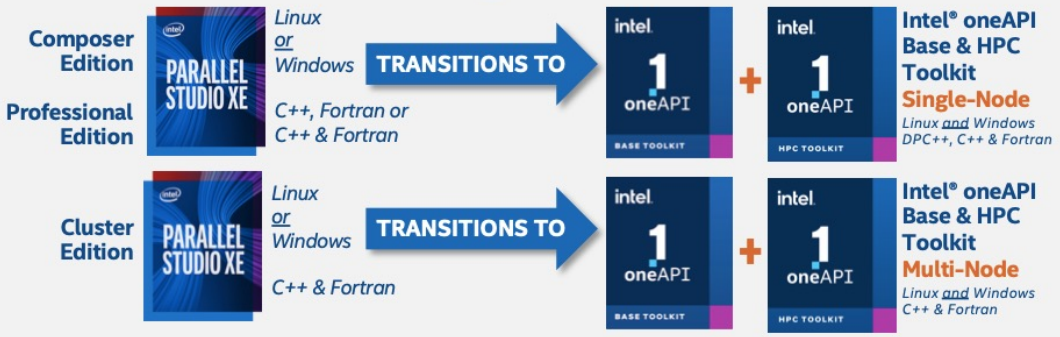

"Licensing" / Usage Enforcement

FLEXlm & Intel EULA
– Named User
- 2 Seat Floating
- 5 Seat Floating

Intel EULA
– Named-user
- 2 Concurrent User, Max 10 Developers Supported
- 5 Concurrent User, Max 25 Developers Supported

Customer Transition to oneAPI

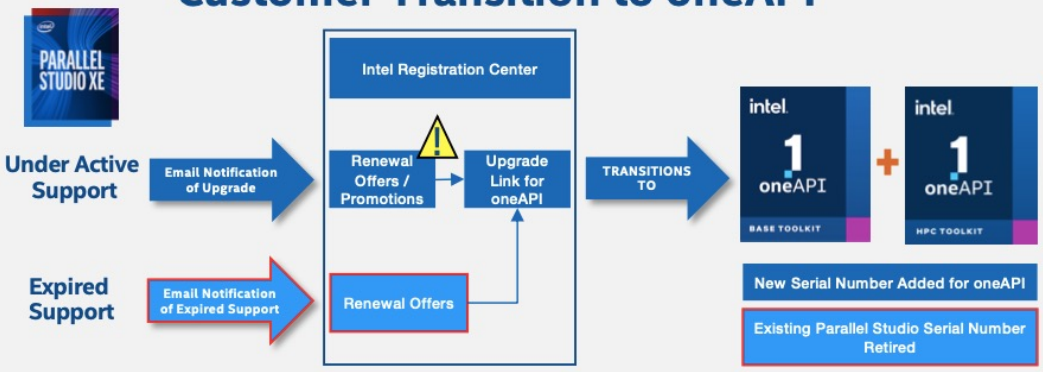

Upgrade Promotion Offers

Intel oneAPI Toolkits Free Availability

Get Started Quickly

Code Samples, Quick-start Guides, Webinars, Training

software.intel.com/oneapi

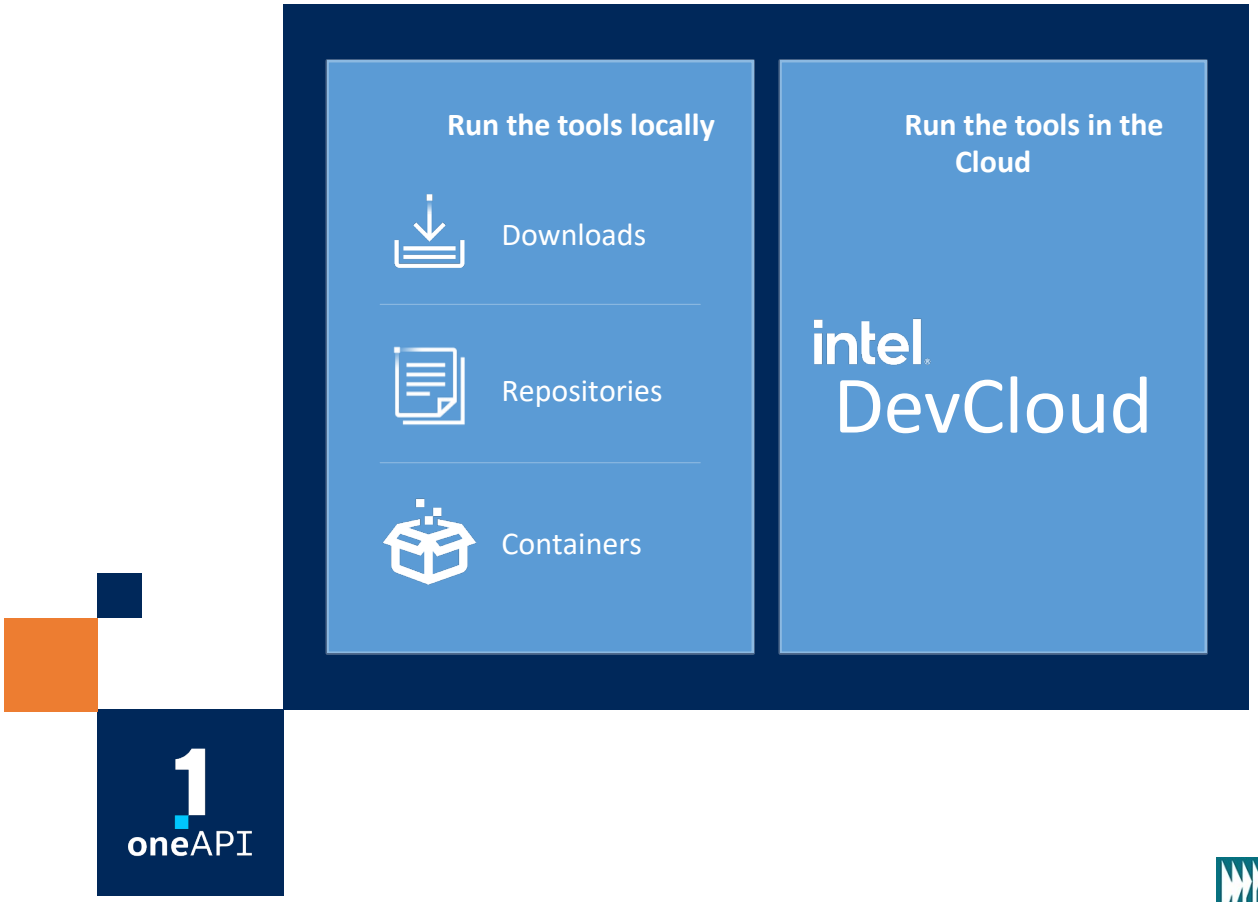

oneAPI Available on: Intel DevCloud

A development sandbox to develop, test and run workloads across a range of Intel CPUs, GPUs, and FPGAs using Intel's oneAPI software.

Get Up & Running In Seconds!

software.intel.com/devcloud/oneapi

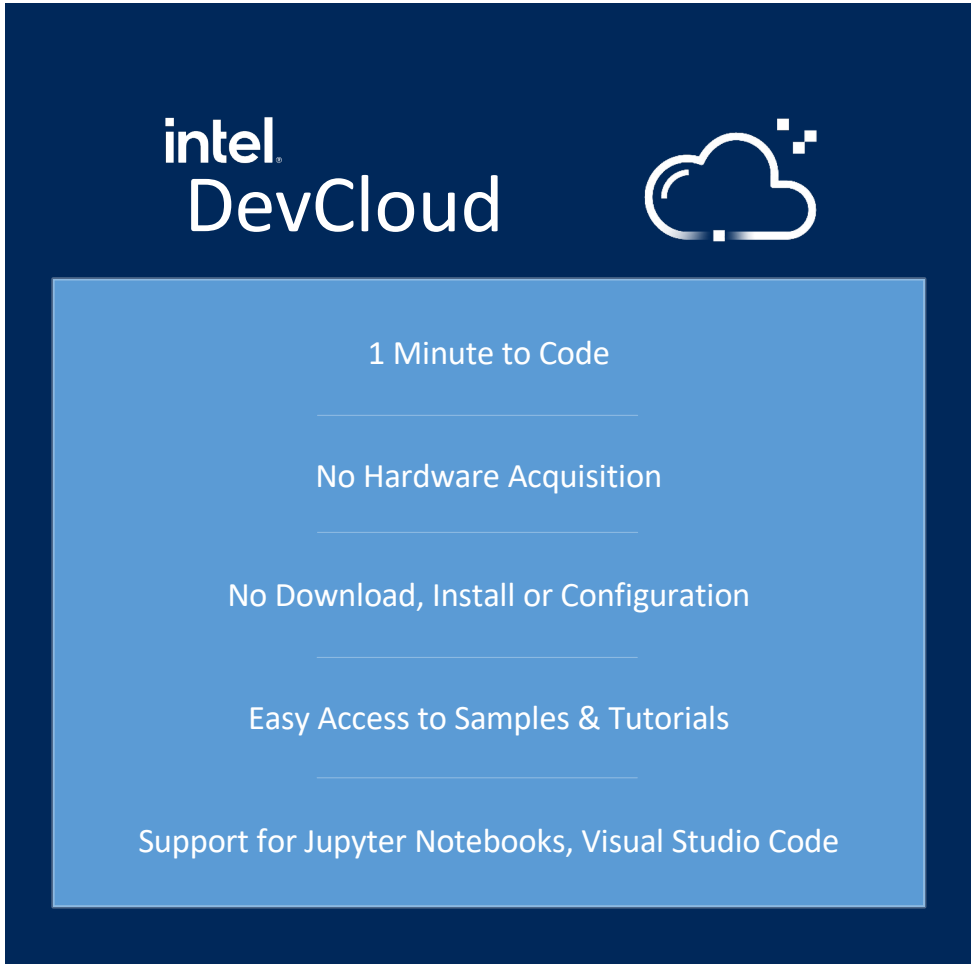

Create an Intel DevCloud Account

software.intel.com/devcloud/oneapi

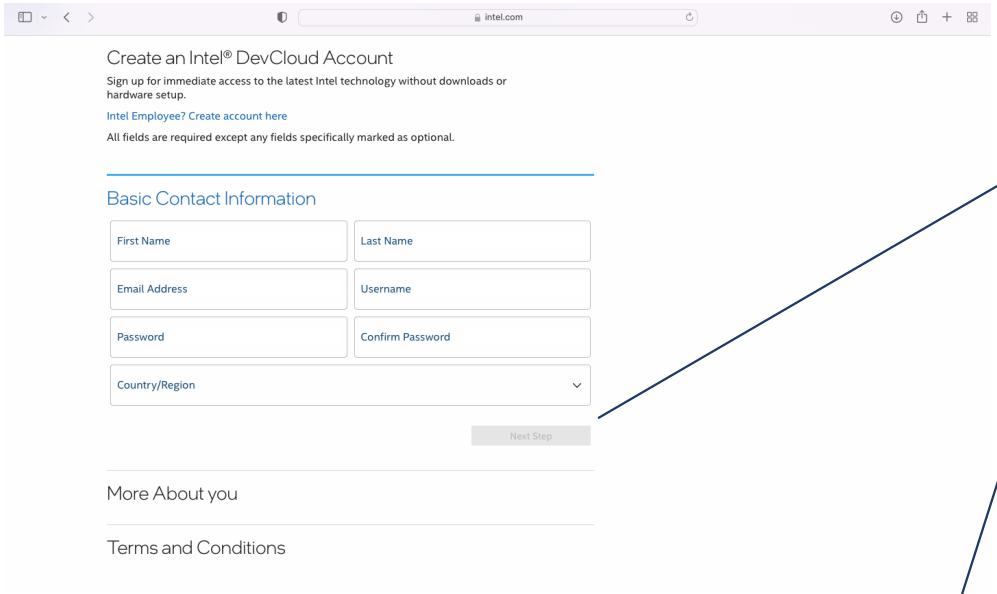

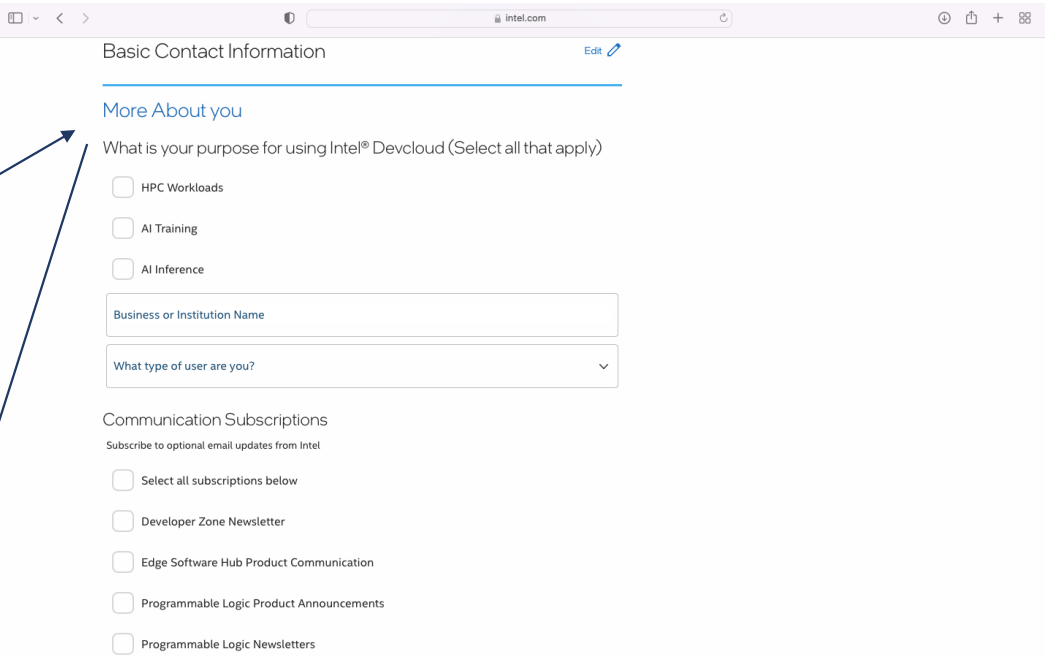

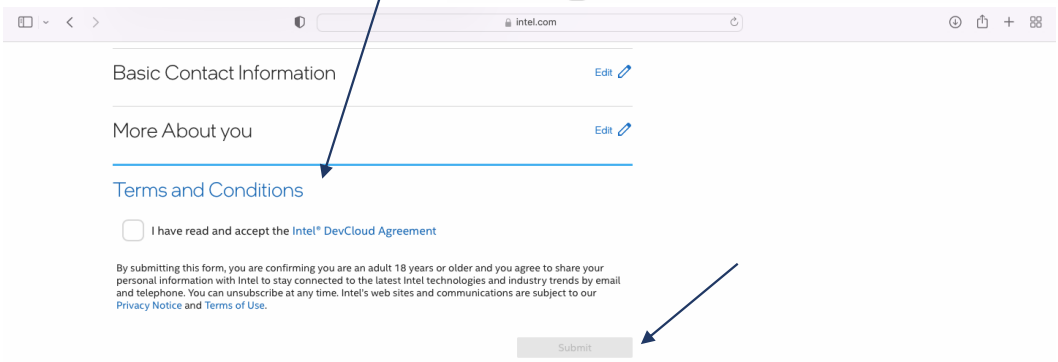

Subject: Intel® DevCloud Account Confirmation - Email Verification

Thank you for registering for an Intel® DevCloud Account.

Please verify your email address by clicking the link below. The link will expire in 5 days.

Verify your email

Your password should be protected as confidential. Your use of the password and Intel's websites are governed by Intel's Terms and Conditions of Use linked from the bottom of each respective site's web pages.

If you have any questions, please contact us.

To manage your profile, including available marketing subscriptions, please visit My Intel.

Please DO NOT reply to this e-mail message. This is an automated response.
To ensure that you continue receiving our e-mails, please add us to your address book.

Intel is committed to protecting your privacy. For more information, please see the Intel Privacy Notice.

intel

Subject: Welcome to Intel® DevCloud

Your Intel® DevCloud account has been activated! You are now able to develop, test, and run your workloads across a range of Intel® CPUs, GPUs, FPGAs and Edge Devices using the latest Intel® software.

Access a variety of tools to get started by visiting https://software.intel.com/devcloud

Your access to Intel DevCloud expires on

You can extend your access within 30 days of expiration. For you privacy, your account and data will be deleted upon expiration, so make sure to backup any data you wish to preserve before then.

Thank you for using Intel® DevCloud.

Access a variety of tools to get started

devcloud.intel.com/oneapi

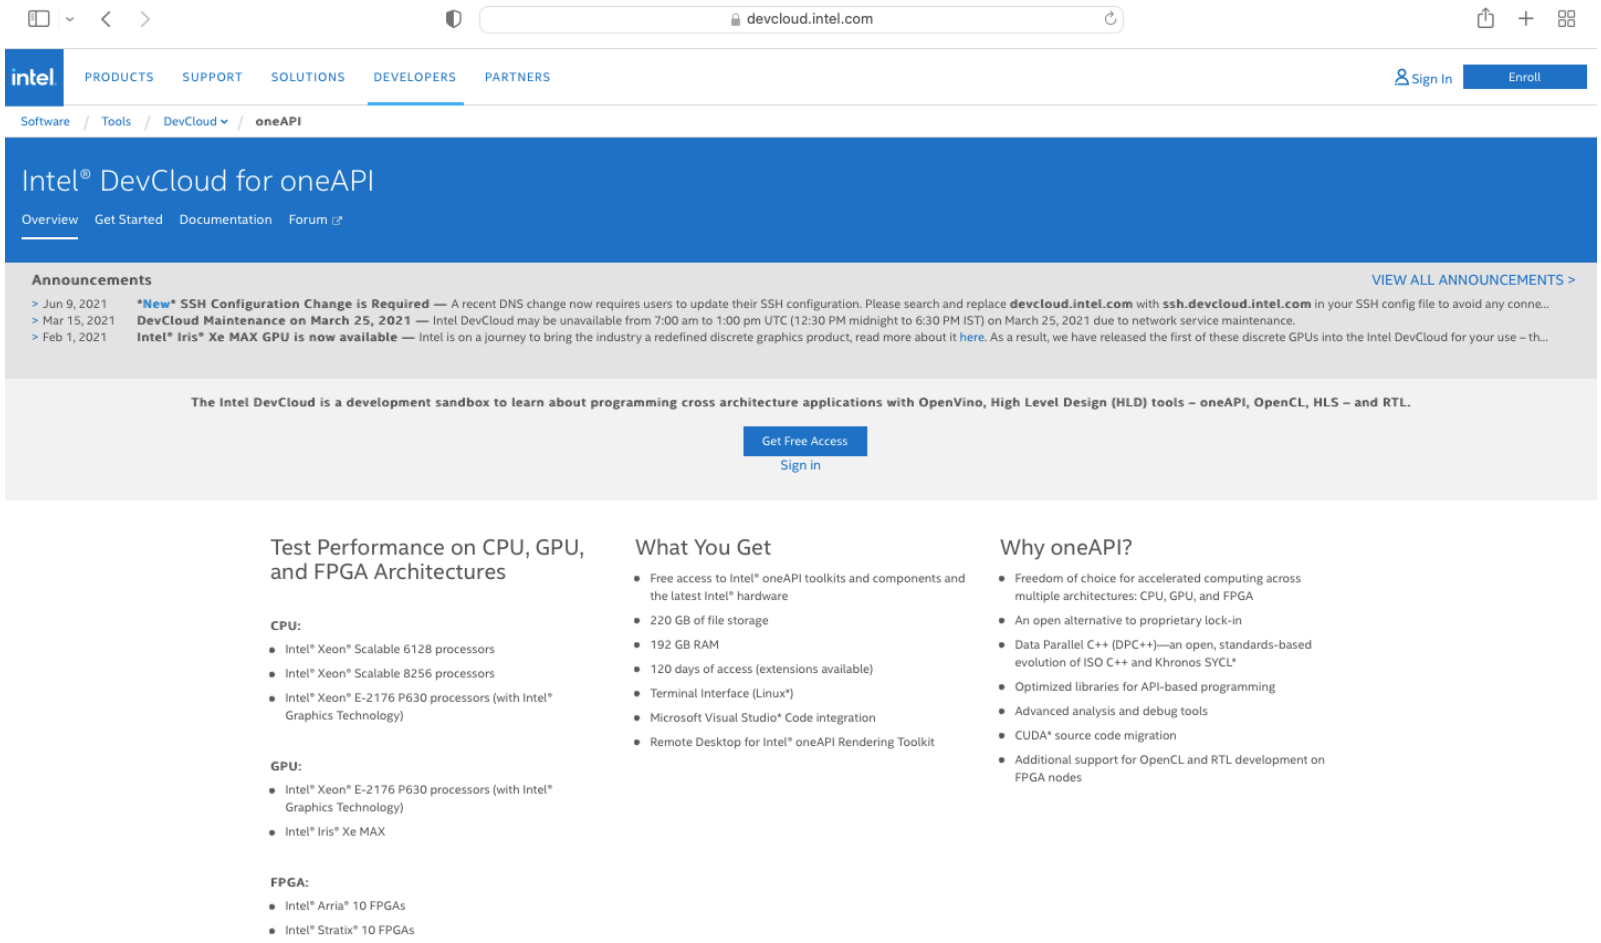

Choose your Connection Method

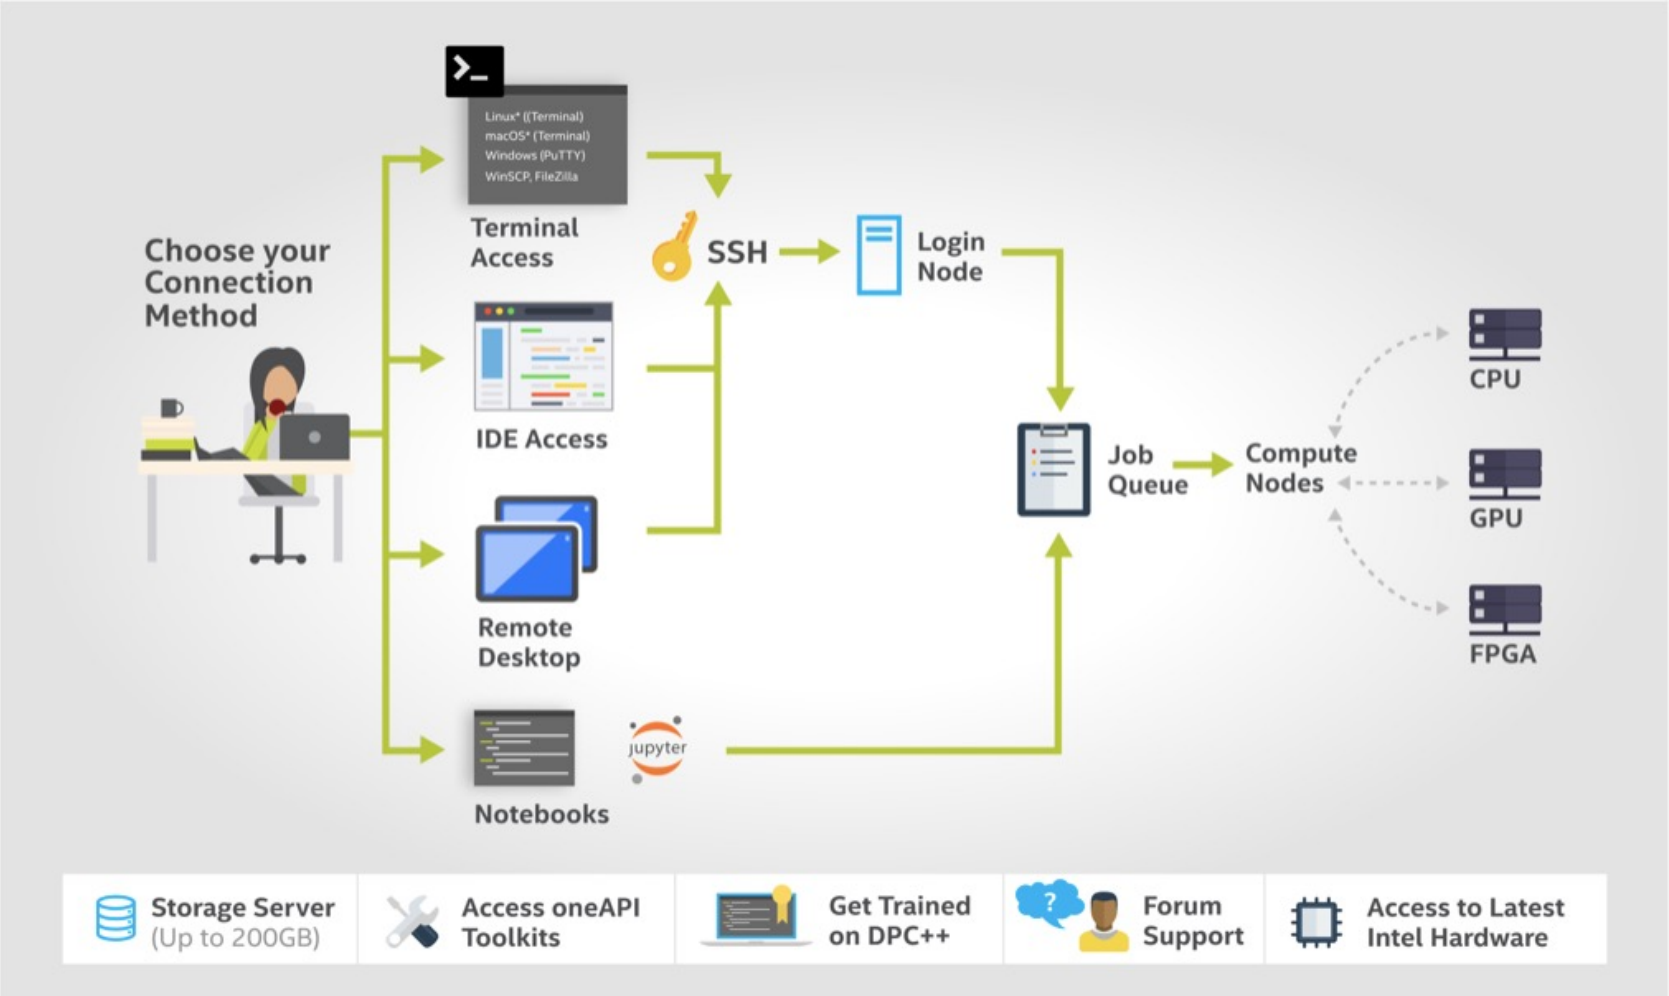

Setup of Intel DevCloud and JupyterLab Environment

Launch Jupyterlab and select Terminal

jupyter.oneapi.devcloud.intel.com

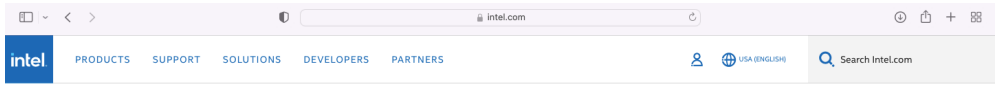

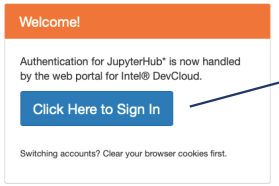

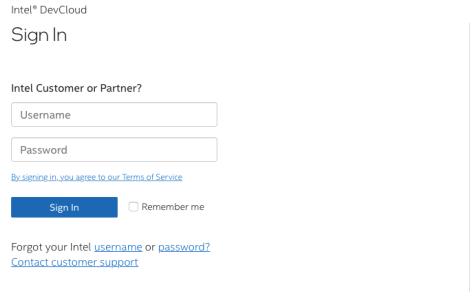

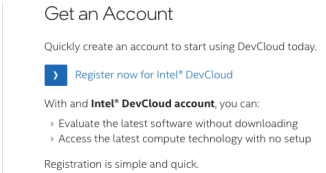

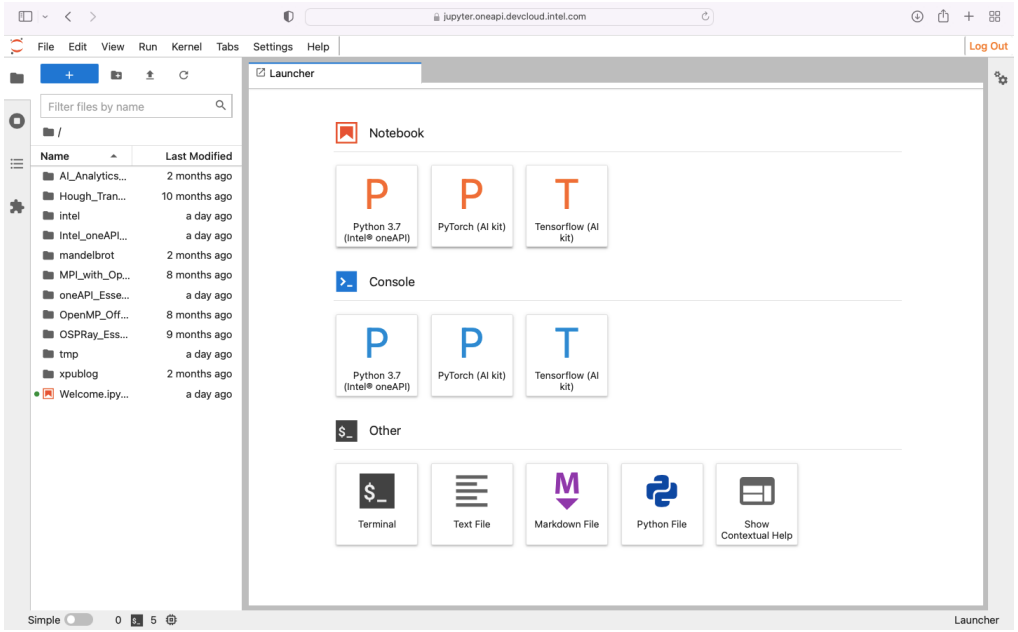

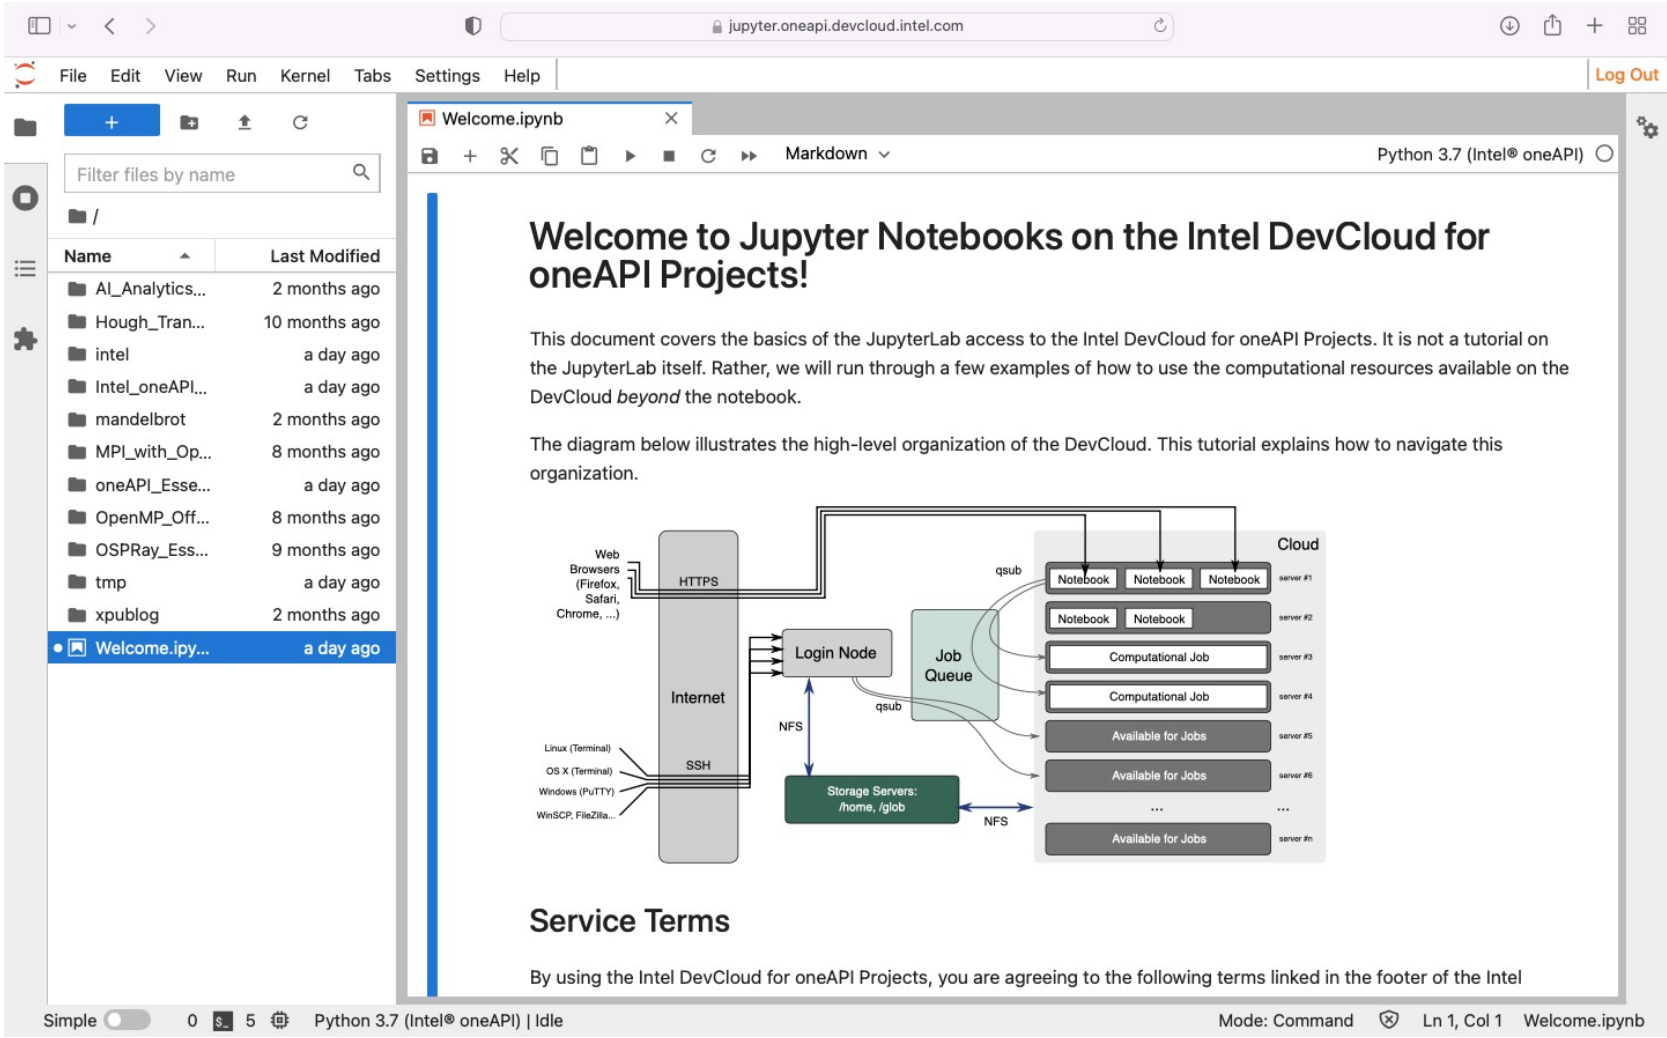
jupyter.oneapi.devcloud.intel.com
File Edit View Run Kernel Tabs Settings Help
Log Out
Filter files by name
/
Name
Last Modified
AI_Analytics... 2 months ago
Hough_Tran... 10 months ago
intel a day ago
Intel_oneAPI... a day ago
mandelbrot 2 months ago
MPI_with_Op... 8 months ago
oneAPI_Esse... a day ago
OpenMP_Off... 8 months ago
OSPRay_Ess... 9 months ago
tmp a day ago
xpublog 2 months ago
Welcome.ipy... a day ago
Welcome.ipynb
Markdown
Python 3.7 (Intel® oneAPI)
Welcome to Jupyter Notebooks on the Intel DevCloud for oneAPI Projects!
This document covers the basics of the JupyterLab access to the Intel DevCloud for oneAPI Projects. It is not a tutorial on the JupyterLab itself. Rather, we will run through a few examples of how to use the computational resources available on the DevCloud beyond the notebook.
The diagram below illustrates the high-level organization of the DevCloud. This tutorial explains how to navigate this organization.
Web Browsers (Firefox, Safari, Chrome, ...)
HTTPS
Internet
Linux (Terminal)
OS X (Terminal)
Windows (PuTTY)
WinSCP, FileZilla...
SSH
Login Node
NFS
qsub
Job Queue
Storage Servers: /home, /glob
Cloud
Notebook
Computational Job
Available for Jobs
Service Terms
By using the Intel DevCloud for oneAPI Projects, you are agreeing to the following terms linked in the footer of the Intel
Simple 0 5 Python 3.7 (Intel® oneAPI) | Idle
Mode: Command Ln 1, Col 1 Welcome.ipynb

Connect with Linux/macOS SSH Client

devcloud.intel.com/oneapi/documentation/connect-with-ssh-linux-macos/

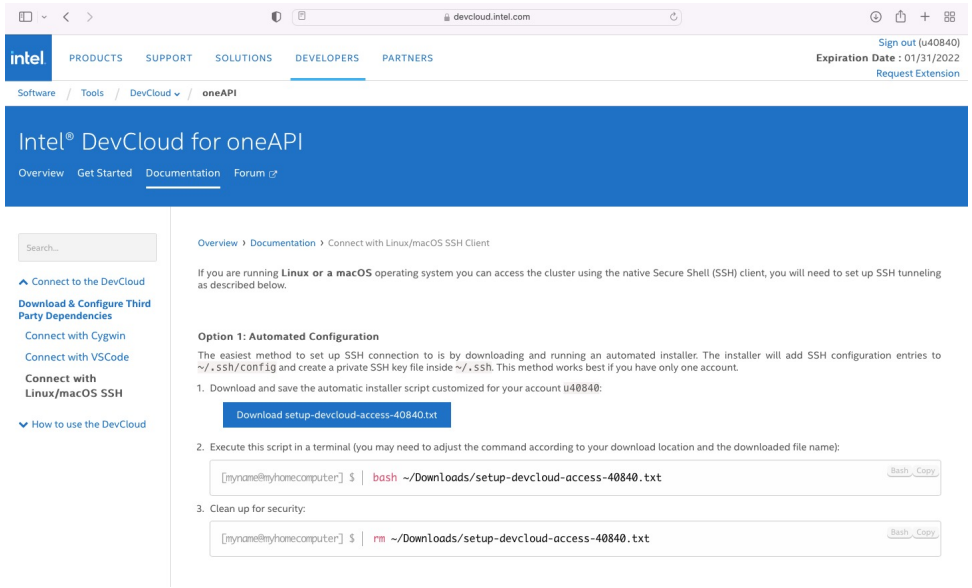

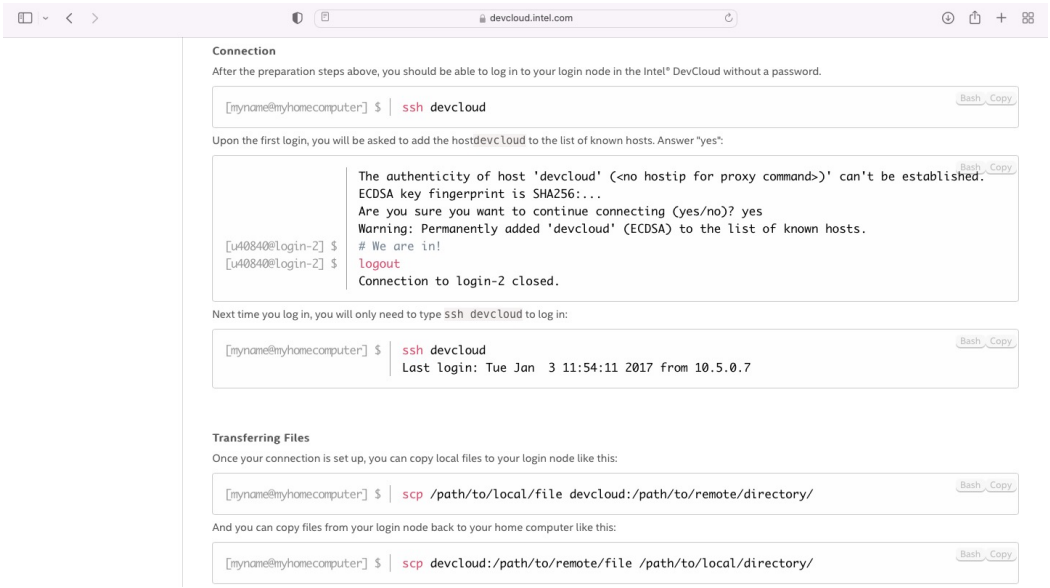

A Complete DPC++ Program

Single source

- Host code and heterogeneous accelerator kernels can be mixed in same source files

Familiar C++

- Library constructs add functionality, such as:

Construct	Purpose
queue	Work targeting
malloc_shared	Data management
parallel_for	Parallelism

Host code

Accelerator device code

Host code

```
#include <CL/sycl.hpp>
constexpr int N=16;
using namespace sycl;
int main() {
  queue q;
  int *data = malloc_shared<int>(N, q);
  q.parallel_for(N, [=](auto i) {
      data[i] = i;
  }).wait();
  for (int i=0; i<N; i++) std::cout << data[i] << "\n";
  free(data, q);
  return 0;
}
```

SYCL Classes

Device

- The **device** class represents the capabilities of the accelerators in a oneAPI system.
- The device class contains member functions for querying information about the device, which is useful for DPC++ programs where multiple devices are created.
- The function get_info gives information about the device:

```
queue q;
device my_device = q.get_device();
std::cout << "Device: " << my_device.get_info<info::device::name>() << std::endl;
```

Device Selector

- The **device_selector** class enables the runtime selection of a particular device to execute kernels based upon user-provided heuristics.
- The following code sample shows use of the standard device selectors (default_selector, cpu_selector, gpu_selector…) and a derived **device_selector**

```
default_selector selector;
// host_selector selector;
// cpu_selector selector;
// gpu_selector selector;
queue q(selector);
std::cout << "Device: " << q.get_device().get_info<info::device::name>() << std::endl;
```

Queue

- A queue **submits command groups** to be executed by the SYCL runtime
- Queue is a mechanism where work is submitted to a device.
- A Queue map to one device and multiple queues can be mapped to the same device.

```
queue q;

q.submit([&](handler& h) {
    // COMMAND GROUP CODE
});
```

Choosing Where Device Kernels Run

Work is submitted to queues

- Each queue is associated with exactly one device (e.g. a specific GPU or FPGA)
- You can:
  - Decide which device a queue is associated with (if you want)
  - Have as many queues as desired for dispatching work in heterogeneous systems

Create queue targeting any device:	queue();

Create queue targeting a pre-configured classes of devices:	queue(cpu_selector{});
	queue(gpu_selector{});
	queue(intel::fpga_selector{});
	queue(accelerator_selector{});
	queue(host_selector{});

Create queue targeting specific device (custom criteria):	class custom_selector : public device_selector {
	int operator()(…… // Any logic you want!
	…
	queue(custom_selector{});

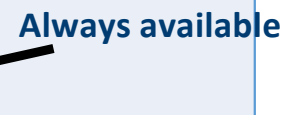

Kernel

- The kernel class encapsulates methods and data for executing code on the device when a command group is instantiated
- Kernel object is not explicitly constructed by the user
- Kernel object is constructed when a kernel dispatch function, such as parallel_for, is called

```
q.submit([&](handler& h) {
  h.parallel_for(range<1>(N), [=](id<1> i) {
    A[i] = B[i] + C[i];
  });
});
```

DPC++ language and runtime

- DPC++ language and runtime consists of a set of C++ classes, templates, and libraries
- **Application scope** and **command group scope**:
  - Code that executes on the host
  - The full capabilities of C++ are available at application and command group scope
- **Kernel scope**:
  - Code that executes on the device.
  - At kernel scope there are limitations in accepted C++

Parallel Kernels

- Parallel Kernel allows multiple instances of an operation to execute in parallel.
- Useful to offload parallel execution of a basic for-loop in which each iteration is completely independent and in any order.
- Parallel kernels are expressed using the parallel_for function

for-loop in CPU application

```
for(int i=0; i < 1024; i++){
    a[i] = b[i] + c[i];
});
```

Offload to accelerator using **parallel_for**

```
h.parallel_for(range<1>(1024), [=](id<1> i){
    A[i] =  B[i] + C[i];
});
```

Basic Parallel Kernels

The functionality of basic parallel kernels is exposed via range, id and item classes

- range class is used to describe the iteration space of parallel execution
- id class is used to index an individual instance of a kernel in a parallel execution
- item class represents an individual instance of a kernel function, exposes additional functions to query properties of the execution range

```
h.parallel_for(range<1>(1024), [=](id<1> idx){
        // CODE THAT RUNS ON DEVICE
});
```

```
h.parallel_for(range<1>(1024), [=](item<1> item){
        auto idx = item.get_id();
        auto R = item.get_range();
        // CODE THAT RUNS ON DEVICE
});
```

ND-Range Kernels

Basic Parallel Kernels are easy way to parallelize a for-loop but does not allow performance optimization at hardware level.

ND-Range kernel is another way to expresses parallelism which enable low level performance tuning by providing access to local memory and mapping executions to compute units on hardware.

- The entire iteration space is divided into smaller groups called work-groups, work-items within a work-group are scheduled on a single compute unit on hardware.
- The grouping of kernel executions into work-groups will allow control of resource usage and load balance work distribution.

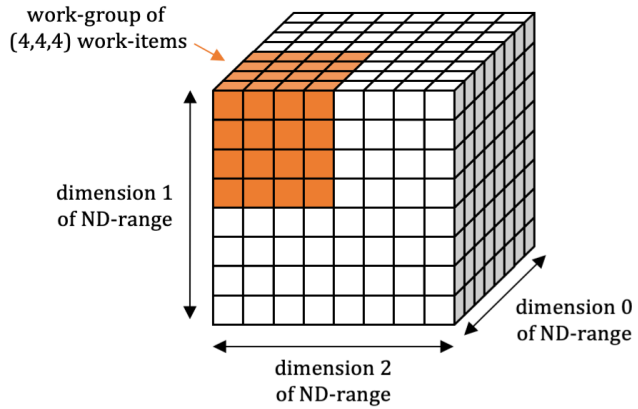

ND-Range

ND-Range Kernels

The functionality of nd_range kernels is exposed via nd_range and nd_item classes

```
h.parallel_for(nd_range<1>(range<1>(1024),range<1>(64)), [=](nd_item<1> item){
    auto idx = item.get_global_id();
    auto local_id = item.get_local_id();
    // CODE THAT RUNS ON DEVICE
});
```

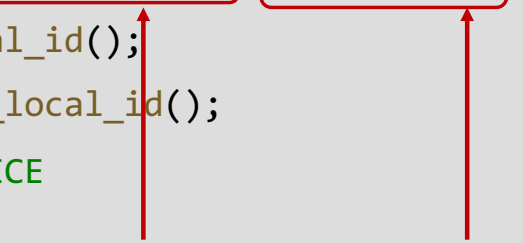

- nd_range class represents a grouped execution range using global execution range and the local execution range of each work-group.
- nd_item class represents an individual instance of a kernel function and allows to query for work-group range and index.

Buffer Memory Model

Buffers: Encapsulate data in a SYCL application

- Across both devices and host!

Accessors: Mechanism to access buffer data

- Create data dependencies in the SYCL graph that order kernel executions

```
queue q;
std::vector<int> v(N, 10);
{
  buffer buf(v);
  q.submit([&](handler& h) {
    accessor a(buf, h , write_only);
    h.parallel_for(N, [=](auto i) { a[i] = i; });
  });
}
for (int i = 0; i < N; i++) std::cout << v[i] << " ";
```

DPC++ Code Anatomy

- oneAPI programs require the include of CL/sycl.hpp
- It is recommended to employ the namespace statement to save typing repeated references into the sycl namespace

```
#include <CL/sycl.hpp>
using namespace sycl;
```

DPC++ Code Anatomy

```
void dpcpp_code(int* a, int* b, int* c) {
  // Setting up a DPC++ device queue
  queue q;
  // Setup buffers for input and output vectors
  buffer buf_a(a, range<1>(N));
  buffer buf_b(b, range<1>(N));
  buffer buf_c(c, range<1>(N));
  //Submit Command group function object to the queue
  q.submit([&](handler &h){
    //Create device accessors to buffers allocated in global memory
    accessor A(buf_a, h, read_only);
    accessor B(buf_b, h, read_only);
    accessor C(buf_c, h, write_only);
    //Specify the device kernel body as a lambda function
    h.parallel_for(range<1>(N), [=](auto i){
      C[i] = A[i] + B[i];
    });
  });
}
```

Step 1: create a device queue (developer can specify a device type via device selector or use default selector)

Step 2: create buffers (represent both host and device memory)

Step 3: submit a command for (asynchronous) execution

Step 4: create buffer accessors to access buffer data on the device

Step 5: send a kernel (lambda) for execution

Step 6: write a kernel

Kernel invocations are executed in parallel

Kernel is invoked for each element of the range

Kernel invocation has access to the invocation id

Done!
The results are copied to vector `c` at `buf_c` buffer destruction

Custom Device Selector

- The following code shows derived **device_selector** that employs a device selector heuristic. The selected device prioritizes a GPU device because the integer rating returned is higher than for CPU or other accelerator.

```
#include <CL/sycl.hpp>
using namespace cl::sycl;
class my_device_selector : public device_selector {
public:
  int operator()(const device& dev) const override {
    int rating = 0;
    if (dev.is_gpu() & (dev.get_info<info::device::name>().find("Intel") != std::string::npos))
      rating = 3;
    else if (dev.is_gpu()) rating = 2;
    else if (dev.is_cpu()) rating = 1;
    return rating;
  };
};
int main() {
  my_device_selector selector;
  queue q(selector);
  std::cout << "Device: " << q.get_device().get_info<info::device::name>() << std::endl;
  return 0;
}
```

Hands-on Coding on Intel DevCloud

Intel VTune Profiler: DPC++ Profiling - Tune for CPU, GPU & FPGA

Analyze Data Parallel C++ (DPC++)

See the lines of DPC++ that consume the most time

Tune for Intel CPUs, GPUs & FPGAs

Optimize for any supported hardware accelerator

Optimize Offload

Tune OpenMP offload performance

Wide Range of Performance Profiles

CPU, GPU, FPGA, threading, memory, cache, storage…

Supports Popular Languages

DPC++, C, C++, Fortran, Python, Go, Java, or a mix

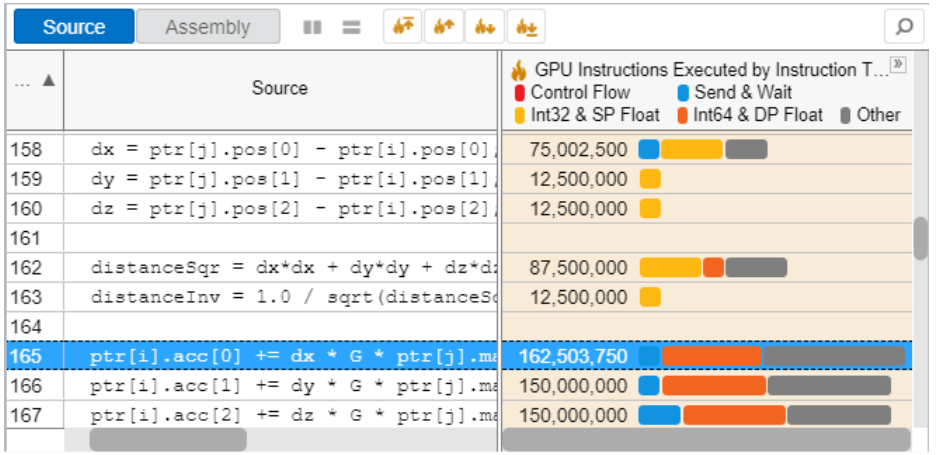

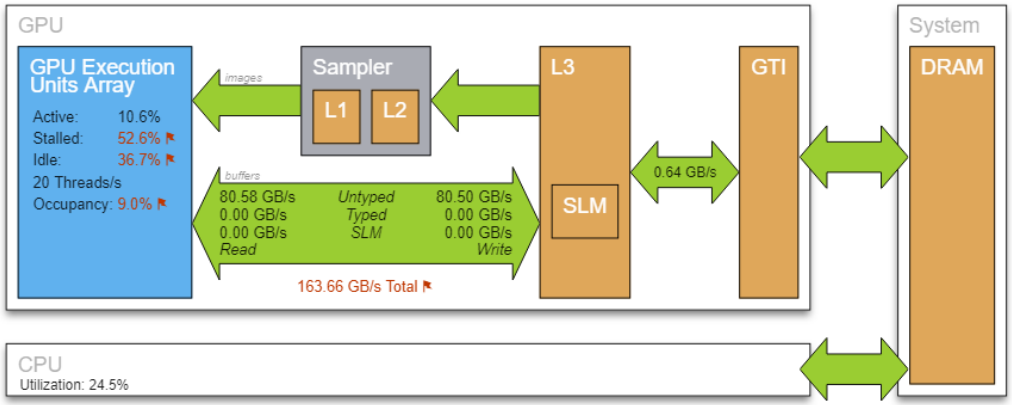

There will still be a need to tune for each architecture.

Hands-on Intel VTune Profiler on Intel DevCloud

Intel DPC++ Compatibility Tool: Minimizes Code Migration Time

- Assists developers migrating code written in CUDA to DPC++ once, generating **human readable** code wherever possible
- ~80-90% of code typically migrates automatically
- Inline comments are provided to help developers finish porting the application

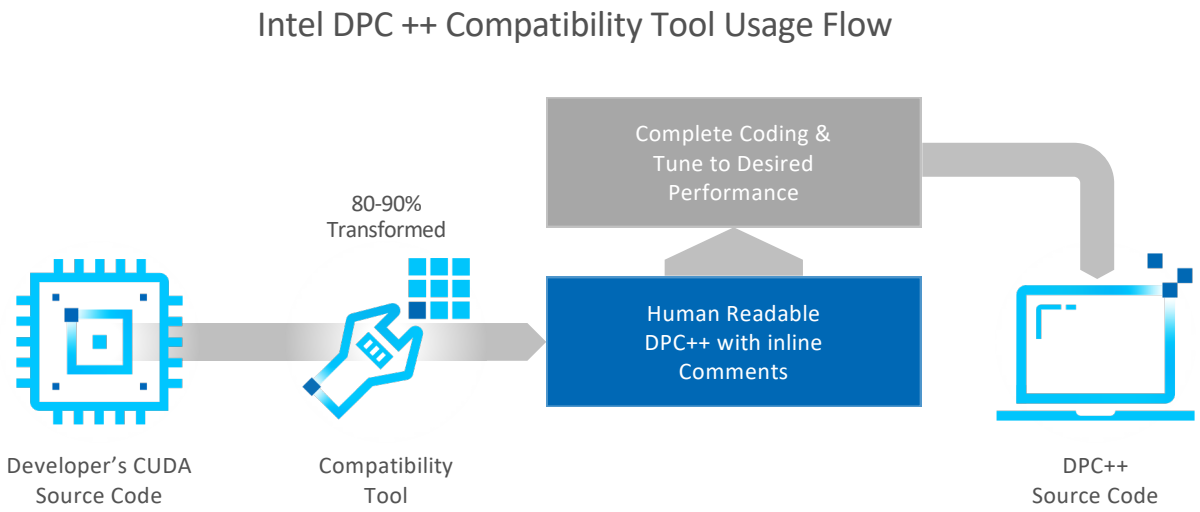

Hands-on Intel DPC++ Compatibility Tool on Intel DevCloud

Thanks!